


# Extending and Updating the Tool Interfaces in MPI: A Request for Feedback

Martin Schulz  
Technische Universität München  
Fakultät für Informatik

Scalable Tools Workshop 2018  
Solitude, UT  
July 2018

With material from

- Marc-Andre Hermanns, JSC
- Kathryn Mohror, LLNL



# Tools Activities in the MPI Forum

Tools WG – Leads:

- Marc-Andre Hermanns, JSC
- Kathryn Mohror, LLNL

Focuses on all aspects of tool interfaces in MPI

- Debugging and performance tools
- Impact on other parts of the standard

Currently under discussion

- **MPI\_T Events – adding callbacks to MPI\_T**
- **QMPI – modernizing PMPI**
- UUIDs for variables and events – easier identification and tracking
- Timers – integers instead of doubles
- Debug interface vs. PMI / PMIx
- What do “Sessions” mean for tools?




# Part 1: MPI\_T Events

## Motivation

- PMPI does not provide access to MPI internal state information
- MPI\_T performance variables only provide aggregated information

Didn't we see the idea of MPI events for tools before? Yes: MPI Peruse

- Access to specific runtime events
- List of point-to-point events defined
- Prototyped, but never standardized



# MPI\_T Events Builds on the Ideas of MPI\_T

Do not mandate specific implementation of MPI functionality

- No requirement to implement specific events

Provide access to MPI implementation-internal information about events

- What happens and when it happens

Notification of events can be immediate or deferred

- Queuing of events can reduce overhead
- It may be impossible to provide immediate notification of some events

Matches the concepts of the existing MPI\_T interface

- Interface
  - A) to query available events (query variables)
  - B) register callbacks (allocate handles)
  - C) read data during callbacks (read variables)

# Complete MPI\_T Events API



Name	Arguments
EVENT TYPE INFORMATION	
MPI_T_event_get_num	<i>int*</i> num_events
MPI_T_event_get_info	<i>int</i> event_index, <i>char*</i> name, <i>int*</i> name_len, <i>int*</i> verbosity, <i>Datatype*</i> array_of_datatypes, <i>MPI_Aint*</i> array_of_displacements, <i>int*</i> num_elements, <i>MPI_T_enum*</i> enum, <i>MPI_Aint*</i> extent, <i>char*</i> description, <i>int*</i> description_len, <i>int*</i> bind
MPI_T_event_get_index	<i>char*</i> name, <i>int*</i> event_index
CALLBACK REGISTRATION MANAGEMENT	
MPI_T_event_handle_alloc	<i>int</i> event_index, <i>void*</i> object_handle, <i>void*</i> user_data, <i>MPI_T_event_cb_function</i> event_cb_function, <i>MPI_T_event_registration*</i> event_registration
MPI_T_event_handle_free	<i>MPI_T_event_registration</i> event_registration, <i>MPI_T_event_free_cb_function</i> free_cb_function
MPI_T_event_set_dropped_handler	<i>MPI_T_event_registration</i> event_registration, <i>MPI_T_event_dropped_cb_function</i> dropped_cb_function
READING EVENT DATA	
MPI_T_event_read	<i>MPI_T_event_instance</i> event, <i>int</i> element_index, <i>void*</i> buffer, <i>int</i> size
MPI_T_event_copy	<i>MPI_T_event_instance</i> event, <i>void*</i> buffer, <i>int</i> size
READING EVENT METADATA	
MPI_T_event_get_wtime	<i>MPI_T_event</i> event, <i>double*</i> event_time
MPI_T_event_get_source	<i>MPI_T_event</i> event, <i>int*</i> source_index
SOURCE HANDLING	
MPI_T_source_get_num	<i>int*</i> num_sources
MPI_T_source_get_info	<i>int</i> source_index, <i>char*</i> name, <i>int*</i> name_len, <i>char*</i> description, <i>int*</i> description_len, <i>MPI_T_source_order*</i> ordering

# Query API



Query  
available  
events  
and their  
semantic  
info

	MPI_T_EVENT_GET_INFO(event_index, name, name_len, verbosity, array_of_datatypes, array_of_displacements, num_datatypes, enumtype, extent, desc, desc_len, bind)		
IN	event_index		index of the event type to be queried; in the range of [0, <i>num_events</i> ) (integer)
OUT	name		buffer to return the string containing the name of the event type (string)
INOUT	name_len		length of the string and/or buffer for <b>name</b> (integer)
OUT	verbosity		verbosity level of this event type (integer)
OUT	array_of_datatypes		array of MPI basic datatypes used to encode the event data (handle)
OUT	array_of_displacements		array of byte displacements of the elements in the event buffer (integer)
INOUT	num_datatypes		length of <b>array_of_datatypes</b> and <b>array_of_displacements</b> arrays (integer)
OUT	enumtype		optional descriptor for enumeration information (handle)
OUT	extent		number of bytes needed for a buffer to copy all data, including padding, encoded in the event type (integer)
OUT	desc		buffer to return the string containing a description of the event type (string)
INOUT	desc_len		length of the string and/or buffer for <b>desc</b> (integer)
OUT	bind		type of MPI object to which an event of this type must be bound (integer)

# Allocating Event Handles and their Callbacks

**Register for events of interest:**

<b>MPI_T_EVENT_HANDLE_ALLOC</b> (event_index, obj_handle, user_data, event_cb_function, handle)		
IN	event_index	index of the event type to be queried between 0 and <i>num_events</i> – 1 (integer)
IN	obj_handle	pointer to a handle of the MPI object to which this event is supposed to be bound (pointer)
IN	user_data	pointer to a user-controlled buffer (pointer)
IN	event_cb_function	pointer to user-defined callback function (pointer)
OUT	handle	allocated handle (handle)

```
typedef void (*MPI_T_event_cb_function)(  
    MPI_T_event event,  
    MPI_T_event_handle handle,  
    MPI_T_cb_safety cb_safety,  
    void *user_data);
```

# Receiving Callbacks



Callbacks for allocated handles are triggered when the corresponding event happens

- Opaque MPI\_T event type can be queried for information
- Type scheme still under discussion

`MPI_T_EVENT_READ(event, element_index, buffer)`

IN	<code>event</code>	event data handle provided to the callback function (handle)
IN	<code>element_index</code>	index into the array of datatypes of the item to be queried (integer)
OUT	<code>buffer</code>	buffer to a memory location to store the item data (pointer)

`MPI_T_EVENT_READ_ALL(event, array_of_buffers)`

IN	<code>event</code>	event data handle provided to the callback function (handle)
OUT	<code>array_of_buffers</code>	array of buffers to a memory locations to store the event data (pointer)

# Special Provisions

Handling of calling safety for callbacks

- Only minimal MPI usage allowed
- Each callback can state the “safety level” at each event instance
- None, Reentrant, thread safe, async signal safe

MPI\_T Events implementations allowed to defer events

- Provide timestamps to match up deferred events

MPI\_T Events implementations allowed to drop events

- Should be the exception, but can be necessary
- Special dropped event handler to indicate dropping to tool

Ordering of events

- Concept of event sources
- Events from the same source are ordered
- Events from different sources can be out of order

# Status: MPI\_T Events

Proposal mostly complete

- [https://github.com/mpiwg-tools/tools-issues/wiki/MPI\\_T-Events](https://github.com/mpiwg-tools/tools-issues/wiki/MPI_T-Events)
- Current proposal text available on request
- “Reading” planned for September MPI Forum meeting

Prototype implementation close to being done

- Based on Open MPI
- Providing Peruse functionality

Publication

- Enabling callback-driven runtime introspection via MPI\_T  
Hermanns, Hjelm, Knobloch, Mohror, Schulz  
To appear in EuroMPI 2018

## Part 2: QMPI

Proposal to redesign the trusted PMPI interface


### Motivation

- Weak symbol intersection is brittle
- Limited to a single tool (unless you use the awesome P<sup>n</sup>MPI)
- Forces tools to be monolithic


### Requirements

- Support multiple concurrent tools in a single process
- Link time or runtime enablement
- Low to no overhead when no tool is attached
- No loss of functionality compared to existing PMPI
  - Basically wrapper functionality
- All language bindings (C, mpif.h, use mpi, use mpif08)
  - Tools can implement functionality in C (in one place) regardless of language
- Integration with MPI thread support

# Basic Scenario Targeted at First



# Basic Scenario Targeted at First



# Basis is Still Basic Wrapping

Each tool implements a set of routines it wraps

- Registered at startup

Tools have independent instances

- Separate storage space
- Created by MPI at/before MPI Event

Each tool instance has the following “available”:

- A functional table with all “PMPI” / follow on routines
- A pointer to store internal information

Wrapping process:

```
Int QMPI_X( <normal parameters>, opaque)
{
    qmpi_x_t pqmpi_x;
    MPI_Table_query("QMPI_X", &pqmpi_x, table);
    ... Do work ...
    err=pqmpi_x(..., opaque);
    ... Do work ...
    return err;
}
```

# Status: QMPI

Concept mostly worked out

- <https://github.com/mpiwg-tools/tools-issues/wiki/Interface-to-Replace-PMPI>
- APIs are being defined
- Working on standards text is coming up soon-ish

Active work on

- Initialization / Bootstrapping
- Opaque information passed through
- Ability to cleanly "loop back" to own layer

Prototype implementation in the works

- As PMPI tool that provides the new interface
- Basic wrapping already possible
- Generalization of the next few months


# Part 3: UUIDs for MPI\_T variables

MPI implementations are free to provide whatever variables make sense for their implementation

- Variables are allowed to change between versions of the library and across HW (analog to performance and control variables)
- Want to provide some stability for tools and keep the freedom for implementations

Organization IDs and variable identifiers registered with MPI Forum

- Allows to identify common variables across MPI implementations
- Allows to keep variables across MPI versions uniquely identifiable



Vendors are allowed to use a "foreign" VendorID for a variable that has the same semantics as the corresponding variable

# Part 4: Timers



Issue 1: Timers only provide double, which requires conversions for some sources

Proposal 1: new general timing routines

Proposal 2: new MPI\_T timers, possibly per source (currently preferred)

MPI\_WTICKS\_ELAPSED()

```
MPI_Count MPI_Wticks_elapsed(void)  
INTEGER(KIND=MPI_COUNT_KIND) MPI_Wticks_elapsed()  
INTEGER(KIND=MPI_COUNT_KIND) MPI_WTICKS_ELAPSED()
```

MPI\_WTICKS\_PER\_SECOND()

```
MPI_Count MPI_Wticks_per_second(void)  
INTEGER(KIND=MPI_COUNT_KIND) MPI_Wticks_per_second()  
INTEGER(KIND=MPI_COUNT_KIND) MPI_WTICKS_PER_SECOND()
```

Issue 2: MPI timing routines cannot be called before MPI\_Init

Proposal: ???

# Summary and Request for Feedback

Currently under discussion

- **MPI\_T Events – adding callbacks to MPI\_T**
- **QMPI – modernizing PMPI**
- UUIDs for variables and events – easier identification and tracking
- Timers – integers instead of doubles
- Debug interface vs. PMI / PMIx
- What do “Sessions” mean for tools?

If you have feedback, please send it to

- Marc-Andre: [m.a.hermanns@fz-juelich.de](mailto:m.a.hermanns@fz-juelich.de)
- Kathryn Mohror: [mohror1@llnl.gov](mailto:mohror1@llnl.gov)
- Martin Schulz: [schulzm@in.tum.de](mailto:schulzm@in.tum.de)

Or join the WG

- TelCons: Thursday at 8am Pacific Time | 5pm CET
- More Information on Github:
- <https://github.com/mpiwg-tools/tools-issues>

